

# MPI Fundamentals

## ONE-SIDED COMMUNICATION ROUTINES

Thanh-Dang Diep  
High Performance Computing Laboratory  
Faculty of Computer Science and Engineering  
Ho Chi Minh City University of Technology

Email: dang@hcmut.edu.vn

August 16, 2016

### 1 Introduction

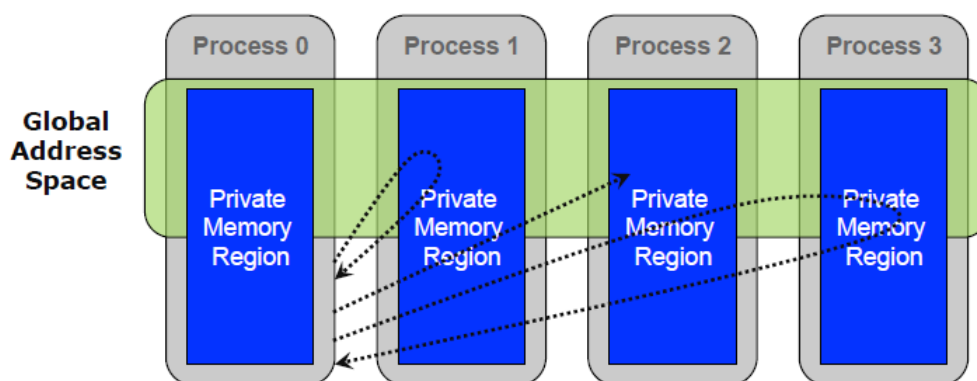


Figure 1: MPI One-sided Communication

The basic idea of one-sided communication models is to decouple data movement with process synchronization:

- Should be able to move data without requiring that the remote process synchronize.
- Each process exposes a part of its memory to other processes.
- Other processes can directly read from or write to this. memory.

## 2 Initialization Routines

### 2.1 MPI\_Win\_create

```
int MPI_Win_create(void *base , MPI_Aint size , int disp_unit ,
                  MPI_Info info , MPI_Comm comm, MPI_Win *win);
```

- Expose a region of memory in an RMA window
  - Only data exposed in a window can be accessed with RMA operations
- Arguments:
  - base** initial address of window (choice)
  - size** size of window in bytes (non-negative integer)
  - disp\_unit** local unit size for displacements, in bytes(positive integer)
  - info** info argument (handle)
  - comm** intra-communicator (handle)
  - win** window object returned by the call (handle)

### 2.2 MPI\_Win\_free

```
int MPI_Win_free(MPI_Win *win);
```

- Free the window object **win** and return a null handle.
- Arguments:
  - win** window object (handle)

**Example 1:**

```

int main(int argc , char **argv)
{
    int *a;
    MPI_Win win;

    MPI_Init(&argc , &argv);

    /* create private memory */
    MPI_Alloc_mem(1000 * sizeof(int) , MPI_INFO_NULL, &a);
    /* use private memory like you normally would */
    a[0] = 1;
    a[1] = 2;

    /* collectively declare memory as remotely accessible */
    MPI_Win_create(a, 1000 * sizeof(int) , sizeof(int) ,
        MPI_INFO_NULL, MPLCOMM_WORLD, &win);

    /* Array 'a' is now accessibly by all processes in
       MPLCOMM_WORLD */

    MPI_Win_free(&win);
    MPI_Free_mem(a);
    MPI_Finalize();
    return 0;
}

```

## 3 Communication Routines

### 3.1 MPI\_Put

```

int MPI_Put(const void *origin_addr , int origin_count ,
            MPI_Datatype origin_datatype , int target_rank ,
            MPI_Aint target_disp , int target_count ,
            MPI_Datatype target_datatype , MPI_Win win);

```

- Move data from origin to target
- Arguments:




Figure 2: Communication using Put

**origin\_addr** initial address of origin buffer (choice)  
**origin\_count** number of entries in origin buffer (non-negative integer)  
**origin\_datatype** datatype of each entry in origin buffer (handle)  
**target\_rank** rank of target (non-negative integer)  
**target\_disp** displacement from start of window to target buffer (non-negative integer)  
**target\_count** number of entries in target buffer (non-negative integer)  
**target\_datatype** datatype of each entry in target buffer (handle)  
**win** window object used for communication (handle)

```
#include <stdio.h>
#include <mpi.h>
#include <stdlib.h>

int main(int argc, char **argv)
{
    int size, rank, localbuf, sharedbuf;
    MPI.Win win;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPLCOMM_WORLD, &size);
    MPI_Comm_rank(MPLCOMM_WORLD, &rank);
```

```

MPI_Win_create(&sharedbuf, 1, sizeof(int), MPI_INFO_NULL
, MPLCOMM_WORLD, &win);

localbuf = rank;

printf("[%d] localbuf = %d\n", rank, localbuf);

MPI_Win_fence(0, win);
MPI_Put(&localbuf, 1, MPI_INT, (rank + 1) % size, 0, 1,
MPI_INT, win);
MPI_Win_fence(0, win);

printf("[%d] sharedbuf = %d\n", rank, sharedbuf);

MPI_Win_free(&win);

MPI_Finalize();
return 0;
}

```

### 3.2 MPI\_Get




Figure 3: Communication using Get

```

int MPI_Get(void *origin_addr , int origin_count ,
            MPI_Datatype origin_datatype , int target_rank ,
            MPI_Aint target_disp , int target_count ,
            MPI_Datatype target_datatype , MPI_Win win );

```

- Move data from target to origin
- Arguments:

**origin\_addr** initial address of origin buffer (choice)

**origin\_count** number of entries in origin buffer (non-negative integer)

**origin\_datatype** datatype of each entry in origin buffer (handle)

**target\_rank** rank of target (non-negative integer)

**target\_disp** displacement from window start to the beginning of target buffer (non-negative integer)

**target\_count** number of entries in target buffer (non-negative integer)

**target\_datatype** datatype of each entry in target buffer (handle)

**win** window object used for communication (handle)

```

#include <stdio.h>
#include <mpi.h>
#include <stdlib.h>

int main(int argc , char **argv)
{
    int size , rank , localbuf , sharedbuf;
    MPI_Win win;

    MPI_Init(&argc , &argv);
    MPI_Comm_size(MPLCOMM_WORLD, &size);
    MPI_Comm_rank(MPLCOMM_WORLD, &rank);

    MPI_Win_create(&sharedbuf , 1 , sizeof(int) , MPI_INFO_NULL
        , MPLCOMM_WORLD, &win);

    sharedbuf = rank;

    printf("[%d] sharedbuf = %d\n" , rank , sharedbuf);
}

```

```

MPI_Win_fence(0, win);
MPI_Get(&localbuf, 1, MPI_INT, (rank + 1) % size, 0, 1,
        MPI_INT, win);
MPI_Win_fence(0, win);

printf("[%d] localbuf = %d\n", rank, localbuf);

MPI_Win_free(&win);

MPI_Finalize();
return 0;
}

```

### 3.3 MPI\_Accumulate




Figure 4: Communication using Accumulate

```

int MPI_Accumulate(const void *origin_addr, int origin_count,
                  MPI_Datatype origin_datatype, int target_rank,
                  MPI_Aint target_disp, int target_count,
                  MPI_Datatype target_datatype, MPI_Op op,
                  MPI_Win win);

```

- Element-wise atomic update operation, similar to a put
  - Reduces origin and target data into target buffer using **op** argument as combiner

- Predefined operations only, no user-defined operations
- Arguments:
  - origin\_addr** initial address of origin buffer (choice)
  - origin\_count** number of entries in origin buffer (non-negative integer)
  - origin\_datatype** datatype of each entry (handle)
  - target\_rank** rank of target (non-negative integer)
  - target\_disp** displacement from start of window to the beginning of target buffer (non-negative integer)
  - target\_count** number of entries in target buffer (non-negative integer)
  - target\_datatype** datatype of each entry in target buffer (handle)
  - op** reduce operation (handle)
  - win** window object (handle)

## 4 Synchronization Routines

**Access epoch:** contain a set of operations issued by an origin process

**Exposure epoch:** enable remote processes to access and/or update a target's window

### 4.1 Fence (Active target)

```
int MPI_Win_fence(int assert , MPI_Win win);
```

- Collective synchronization model
- Starts and ends access and exposure epochs on all processes in the window
- All processes in group of "win" do an MPI\_WIN\_FENCE to open an epoch
- Everyone can issue PUT/GET operations to read/write data
- Everyone does an MPI\_WIN\_FENCE to close the epoch
- All operations complete at the second fence synchronization






Figure 5: Synchronization using Fence

- Arguments:

**assert** program assertion (integer)

**win** window object (handle)

## 4.2 Post-start-complete-wait (Generalized active target)

### 4.2.1 MPI\_Win\_start

```
int MPI_Win_start(MPI_Group group, int assert, MPI_Win win);
```

- Origin opens access epoch
- Arguments:




Figure 6: Synchronization using PSCW

**group** group of target processes (handle)  
**assert** program assertion (integer)  
**win** window object (handle)

#### 4.2.2 MPI\_Win\_complete

```
int MPI_Win_complete(MPI_Win win);
```

- Origin closes access epoch
- Arguments:  
**win** window object (handle)

#### 4.2.3 MPI\_Win\_post

```
int MPI_Win_post(MPI_Group group, int assert, MPI_Win win);
```

- Target open expose epoch

- Arguments:
  - group** group of origin processes (handle)
  - assert** program assertion (integer)
  - win** window object (handle)

#### 4.2.4 MPI\_Win\_wait

```
int MPI_Win_wait(MPI_Win win);
```

- Target closes exposure epoch
- Arguments:
  - win** window object (handle)

### 4.3 Lock/unlock (Passive Target)

#### 4.3.1 MPI\_Win\_lock

```
int MPI_Win_lock(int lock_type , int rank , int assert , MPI_Win win);
```

- Begin an RMA access epoch at the target process
- Arguments:
  - lock\_type** either MPI\_LOCK\_EXCLUSIVE or MPI\_LOCK\_SHARED (state)
  - rank** rank of locked window (non-negative integer)
  - assert** program assertion (integer)
  - win** window object (handle)

#### 4.3.2 MPI\_Win\_unlock

```
int MPI_Win_unlock(int rank , MPI_Win win);
```

- Complete an RMA access epoch at the target process
- Arguments:
  - rank** rank of window (non-negative integer)
  - win** window object (handle)

## Exercises

1. Write a program that prints out prime numbers in the first one billion of positive integers.
2. Write a program that sums up prime numbers in the first one billion of positive integers.
3. Write a program that computes the value of  $\pi$  using Monte Carlo simulation. The program samples points inside the rectangle delimited by  $(0,0)$  and  $(1,1)$  and counts how many of these are within a circle with a radius of 1. The ratio between the number of points inside the circle and the total number of samples is  $\pi/4$ .
4. Write a program multiplying two square matrices whose sizes of each are  $1000 \times 1000$  and  $10000 \times 10000$ .